



Partitions and QoS

Partitions and QoS

- Daniel Trahan
 - *Email:* Daniel.Trahan@Colorado.edu
 - *RC Homepage:* <https://www.colorado.edu/rc>
 - *RC Email:* rc-help@colorado.edu
-
- Slides available for download at:
 - https://github.com/ResearchComputing/Partition_and_QoS_Fall_2020

Outline

- What is a Partition
- Summit Partitions
 - Shas
 - Sgpu
 - Smem
 - Others
- Summit QoS
- Using Partitions and QoS in a Job

Quick note

- Clusters come in all shapes and sizes, so much of the information here may not apply to other systems. HPC is a very diverse landscape so make sure you check with the system administrators of whichever cluster you are using.

What is a Partition

- Research Computing offers a variety of different node types for users to utilize.
- A **partition** is a grouping of node types that offer specialized features that are not found in our most common node.
 - Each node type offers different resource limits.
- By default the partition is “shas” is selected and should be used for most jobs.
- For a full guide of Partitions, checkout our docs!
 - <https://curc.readthedocs.io/en/latest/running-jobs/job-resources.html#partitions>

The shas partition

- Summit's most common partition for users.
 - shas = Summit Haswell Node
 - Selected by default if no partition is provided
- Node Details
 - Node Count: 380 Nodes
 - Core Count: 24 cores
 - Memory limit: 4.84 GB per Core (Max: 116 GB)
- Queue wait: Usually very active but little wait if Allocation is not heavily used.

The sgpu partition

- Summit's specialized GPU partition.
- Each node is equipped with a 2 Tegra K80 Nvidia GPUs
- Not specialized for Machine learning!
- Node Details:
 - Node count: 10
 - Core count: 24
 - Memory: 4.84 GB per core
- GPU Details:
 - 4992 CUDA Cores
 - 24 GB of Graphical RAM
- Queue wait: Can vary substantially depending on time of year.

The smem partition

- Summits specialized High Memory resource
- Very useful for applications that demand high memory resource usage with no cross node communication.
- Max wall time is 7 Days by default.
- Node Details:
 - Node count: 5
 - Core count: 48
 - Memory: 42.7 GB per core (Max 2 TB)
- Queue wait: Usually congested, one of Summit's most demanded resources.

Partition Variants

- The Shas partition has several variants that we make available for users wanting for faster Queue times. These partitions have specialized limits of use:
 - shas-testing: Special high priority partition for shas testing jobs
 - shas-interactive: Special high priority partition for shas interactive jobs
 - sgpu-testing: Special high priority partition for sgpu testing jobs
 - sknl-testing: Special high priority partition for smem testing jobs
- Testing Partitions: Max 30 min runtime @ 24 cores
- Interactive Partitions: Max 4 hour runtime @ 1 core

Other types of Partitions

- Summit Condo Buy-Ins: ssky and ssky-preemptable
 - Summit Skylake nodes
 - Node Count: 5 nodes and 15 nodes on ssky and ssky-preemptable respectively.
 - Core Count: 24 cores per node.
 - Memory: 7.68 GB per core.
- Summit KNL Nodes
 - Summit Phi (KNL) nodes
 - GPU in a CPU: allows high cores per node.
 - Performance was not as good as expected and Intel discontinued.
 - Still around for your use!

What is a QoS

- Summit also leverages a mechanism called QoS to change certain limitations imposed on your Jobs.
- **QoS** or **Quality of Service** is a field that constrains or modifies certain parameters of your job script.
 - Changes Priority of jobs so you may wait shorter or longer depending on the application.
 - Can also change limitation on how many nodes can be run on the system
- Do not set a QoS unless you need the QoS modifications.
- Ex: Biggest limitation
 - By default Summit jobs cannot run longer than 24 hours.
 - Fix: The long QoS!

The Normal QoS

- Summit's base QoS set by default.
- QoS Parameters:
 - Sets a max walltime on all partitions (except smem) to 24 Hours.
 - 1000 Jobs can be submitted per user.
 - Up to 256 Nodes can be requested per Job.

The Long QoS

- Specialized QoS allowing for longer runs on Summit
- QoS Parameters:
 - Sets a max walltime on shas, sknl, and ssy to 7 Days
 - Maximum of 200 Jobs can be submitted at a time
 - Only 22 nodes maximum can be requested per job. Maximum of 40 Nodes total throughout Summit.
- Generally try to avoid this! Queue times can be quite substantial so keep your jobs smaller to avoid the wait.

Using QoS and Partition in Jobs

- To utilize a Summit Partition or QoS, simply add the following sbatch directives to your job scripts:

```
#SBATCH --partition=<desired-partition>  
#SBATCH --qos=<desired-qos>
```

- On interactive jobs, this can be set as such:

```
sinteractive --partition=<desired-partition \\  
--qos=<desired-qos>
```

Reminder: Never set QoS unless you wish to utilize different resource limits for your desired partition.

Example 1:

```
#!/bin/bash

#SBATCH --nodes=1
#SBATCH --ntasks=24
#SBATCH --time=04:00:00
#SBATCH --partition=shas
#SBATCH --job-name=gpu-job-ex
#SBATCH --output=gpu.%j.out

# Load Modules
module load cuda/10.1

# Run Application
./GPUApp
```

Example 2:

```
#!/bin/bash

#SBATCH --nodes=5
#SBATCH --ntasks=120
#SBATCH --time=48:00:00
#SBATCH --partition=shas
#SBATCH --qos=long
#SBATCH --output=long.%j.out

# Load Modules
module load intel impi

# Run Application
./longapp
```

Some Common issues

- **Problem:** My shas-interactive or shas-testing jobs won't run!
- **Solution:** Usually this is because you have reached the job limits of that partition. Since those two partitions have single job limits, try checking to see if you have an existing job running.

- **Problem:** I'm trying to run a high memory node job for 7 days but it taking forever!
- **Solution:** You might have set QoS to long in your job script. Remember high memory node jobs allow for 7 day run time by default.

Questions?

Thank you!

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact information: rc-help@Colorado.edu
- Slides:
https://github.com/ResearchComputing/Partition_and_QoS_Fall_2020
- Documentation:
<https://curc.readthedocs.io/en/latest/running-jobs/job-resources.html>